

NVIDIA-Certified Associate Generative AI LLMs Master Cheat Sheet

Organized quick-review glossary for exam recall, scenario mapping, and final revision.

Prepared by: Yasir Alharbi

Cybersecurity and AI enthusiast

LinkedIn: www.linkedin.com/in/Yasir-T-Alharbi

How to use: scan domains first, then use decision tree and scenario keywords for exam questions.

Main Sections

- Exam Domains
- Repeated Exam Patterns
- AI Subject Terms Glossary
- Python Tools and Libraries
- Exam Decision Tree
- Common Exam Traps
- Formula and Metric Mini Sheet
- Scenario Keywords

Consolidated from the provided PDFs and domain screenshots. Definitions are intentionally short for fast exam review.

Exam Domains

Core ML & AI Fundamentals

Weight: 30%

Concepts

Artificial Intelligence (AI): Systems that perform tasks requiring human intelligence.

Machine Learning (ML): Algorithms learn patterns from data.

Deep Learning (DL): Multi-layer neural networks learn representations automatically.

Generative AI: AI that creates text, images, code, audio, or data.

Foundation Model: Large pretrained model adaptable to many tasks.

Large Language Model (LLM): Transformer model trained on large text corpora.

Supervised Learning: Learns from labeled input-output examples.

Unsupervised Learning: Finds patterns in unlabeled data.

Semi-supervised Learning: Uses small labeled data plus large unlabeled data.

Self-supervised Learning: Creates labels from raw data itself.

Reinforcement Learning: Learns actions using rewards.

Transfer Learning: Reuses pretrained knowledge for a new task.

Feature Extraction: Freeze base model; train only final layers.

Fine-tuning: Continue training a pretrained model on task data.

Feature Engineering: Create useful model inputs from raw data.

Representation Learning: Model learns useful internal features.

Embedding: Dense vector representing meaning or features.

Text Embedding: Vector representation of text semantics.

Word2Vec: Learns word embeddings from text context.

CBOW: Predicts target word from surrounding context.

Skip-gram: Predicts context words from a target word.

Overfitting: Performs well on training data, poorly on unseen data.

Underfitting: Model too simple to learn patterns.

Regularization: Reduces overfitting.

Dropout: Randomly disables neurons during training.

L1 Regularization: Encourages sparse weights.

L2 Regularization: Penalizes large weights.

Loss Function: Measures prediction error.

Gradient Descent: Updates weights to reduce loss.

Learning Rate: Step size for weight updates.

Backpropagation: Computes gradients through network layers.

Epoch: One full pass through training data.

Batch Size: Samples processed per training step.

Hyperparameter: Config set before training.

Parameter: Learned model weight.

Inference: Running a trained model to produce output.

Classification: Predicts categories.

Regression: Predicts continuous values.

Clustering: Groups similar unlabeled data.

Anomaly Detection: Finds unusual data points.

Association Rule Mining: Finds item co-occurrence rules.

Time Series Forecasting: Predicts future ordered values.

Data Augmentation: Creates varied training examples.

Normalization: Scales values for stable training.

Standardization: Centers/scales features by mean and variance.

Padding: Adds tokens/pixels to standardize size.

Masking: Prevents attention to irrelevant positions.

Architectures

Dense Network: Fully connected neural network.

CNN: Learns spatial features from images.

Convolution: Sliding filter extracts local patterns.

Kernel/Filter: Small matrix detecting image features.

Stride: Step size of convolution movement.

Pooling: Reduces spatial size.

Max Pooling: Keeps strongest local activation.

Average Pooling: Averages local activations.

RNN: Processes sequential data step by step.

LSTM: RNN variant for long dependencies.

GRU: Simpler gated RNN variant.

Transformer: Attention-based architecture for sequences.

Encoder-only Model: Reads input for understanding tasks.

Decoder-only Model: Generates text autoregressively.

Encoder-decoder Model: Converts input sequence to output sequence.

BERT: Bidirectional encoder model.

GPT: Decoder-only autoregressive model.

T5: Encoder-decoder text-to-text model.

BART: Encoder-decoder generation model.

Vision Transformer (ViT): Treats image patches like tokens.

Autoencoder: Compresses and reconstructs data.

VAE: Generative autoencoder with latent distribution.

GAN: Generator and discriminator compete.

Diffusion Model: Learns to denoise step by step.

CLIP: Aligns image and text embeddings.

GNN: Learns from graph nodes and edges.

Multimodal LLM: Processes multiple data types together.

Transformer Concepts

Self-attention: Tokens attend to other tokens in same sequence.

Cross-attention: Decoder attends to encoder outputs.

Multi-head Attention: Multiple attention views in parallel.

Query (Q): Vector used to search attention targets.

Key (K): Vector used for attention matching.

Value (V): Vector containing retrieved information.

Scaled Dot-product Attention: Attention using QK scores scaled by $\sqrt{d_k}$.

$\sqrt{d_k}$ Scaling: Prevents large dot products destabilizing softmax.

Positional Encoding: Adds token order information.

Sinusoidal Encoding: Original fixed transformer position signal.

RoPE: Rotary position embedding for long contexts.

Causal Mask: Prevents decoder from seeing future tokens.

Padding Mask: Ignores padded tokens.

Feed-forward Network (FFN): Per-token nonlinear transformation.

Residual Connection: Adds input back to layer output.

LayerNorm: Normalizes activations inside transformer layers.

RMSNorm: Efficient LayerNorm alternative.

Context Window: Maximum tokens a model can process.

KV Cache: Stores attention keys/values for faster decoding.

Grouped-query Attention: Shares KV heads to reduce inference memory.

Sliding Window Attention: Attends to local token windows.

FlashAttention: Faster memory-efficient attention.

Mixture of Experts (MoE): Activates only selected expert submodels.

Applications

RAG Systems: Retrieve knowledge before generation.

Chatbots: Conversational LLM applications.
Summarizers: Shorten text while preserving meaning.
Translation: Convert text between languages.
Sentiment Analysis: Classify emotional tone.
NER: Detect named entities in text.
Token Classification: Label each token.
Text Classification: Label whole text.
Question Answering: Answer from context or generation.
Extractive QA: Selects answer span from passage.
Generative QA: Produces free-form answer.
Image Classification: Predicts image class.
Object Detection: Finds objects in images.
Speech Recognition: Converts audio to text.
Fraud Detection: Flags suspicious behavior.
Customer Churn Prediction: Predicts customer loss.
Recommendation Systems: Suggest items/users/content.
Drug Discovery: Uses models for molecules/proteins.

Frameworks

PyTorch: Dynamic deep learning framework.
TensorFlow: Production ML framework.
Hugging Face Transformers: Library for transformer models.
spaCy: Industrial NLP processing library.
NumPy: Python numeric computing package.
Pandas: CPU DataFrame analysis library.
scikit-learn: Classical ML toolkit.
Matplotlib/Plotly: Visualization libraries.
Jupyter: Interactive notebooks.

Data Analysis Techniques

Weight: 14%

Techniques

Data Mining: Extracts patterns from large data.
Data Visualization: Shows insights with charts.
Statistical Metrics: Quantify model/data behavior.
Trend Analysis: Finds patterns over time.
Performance Comparison: Compares models using metrics.

Exploratory Data Analysis: Understands distributions and relationships.
Descriptive Statistics: Summarizes data.
Correlation Analysis: Measures relationship between variables.
Distribution Analysis: Studies data shape/spread.
Outlier Detection: Finds abnormal values.
Dimensionality Reduction: Compresses features while preserving structure.
PCA: Linear dimensionality reduction.
Nearest Neighbor Search: Finds closest vectors/items.
Cosine Similarity: Compares vector orientation.
Euclidean Distance: Measures straight-line distance.
Clustering Analysis: Groups similar records.
K-means: Clusters around centroids.
Hierarchical Clustering: Builds nested cluster tree.
Spectral Clustering: Uses graph structure for clustering.
Graph Analytics: Analyzes nodes and edges.
Knowledge Graph: Structured entity-relation graph.

Processes

Data Collection: Gather raw data.
Data Cleaning: Fix missing, duplicate, inconsistent data.
Deduplication: Removes repeated records.
Filtering: Removes low-quality or irrelevant data.
Data Transformation: Convert data into model-ready form.
Feature Scaling: Normalize numeric ranges.
Categorical Encoding: Convert categories to numbers.
Tokenization: Break text into tokens.
Chunking: Split documents for retrieval/context limits.
Data Labeling: Assign ground-truth labels.
Data Versioning: Track dataset changes.
Dataset Provenance: Track data origin and usage rights.
Train/Validation/Test Split: Separate training, tuning, final testing.
Insight Generation: Turn analysis into action.

Tools

RAPIDS: GPU-accelerated data science suite.
cuDF: GPU DataFrame library; pandas-like.
cuML: GPU ML algorithms.

cuGraph: GPU graph analytics.
Dask-cuDF: Distributed multi-GPU DataFrames.
NVTabular: GPU feature engineering for tabular data.
CUDA-X: NVIDIA accelerated software libraries.
NVIDIA DALI: GPU data loading/preprocessing.
TAO Toolkit: Simplifies training and augmentation.
Pandas: CPU data manipulation.
Polars: Fast DataFrame library.
Dask: Distributed Python data processing.

Experimentation & Evaluation

Weight: 22%

Methods

A/B Testing: Compare versions on real traffic.
Benchmarking: Compare against standard baselines.
Cross-validation: Evaluate across multiple data splits.
K-fold Validation: Split data into k train/test folds.
Hold-out Testing: Use separate unseen test set.
Temporal Validation: Test using future time data.
Statistical Testing: Checks significance of results.
Human Evaluation: Humans judge quality.
LLM-as-a-Judge: LLM scores outputs.
Error Analysis: Systematically inspect failures.
Ablation Study: Remove components to measure impact.
Stress Testing: Test under heavy/extreme conditions.
Adversarial Testing: Test with hostile edge cases.
Red Teaming: Attack model to find risks.
Canary Deployment: Release to small user subset.

Metrics

Accuracy: Fraction correct.
Precision: Correct positives among predicted positives.
Recall: Correct positives among actual positives.
F1 Score: Harmonic mean of precision and recall.
Confusion Matrix: Counts prediction categories.
Specificity: Correct negatives among actual negatives.
AUC-ROC: Ranking quality across thresholds.

MAE: Average absolute regression error.
MSE: Average squared regression error.
RMSE: Square root of MSE.
R-squared: Regression variance explained.
Cross-entropy Loss: Classification probability loss.
BLEU: N-gram overlap for translation.
ROUGE: Overlap metric for summaries.
METEOR: Translation quality metric.
Perplexity: Language model uncertainty.
Groundedness: Output supported by sources.
Faithfulness: Response matches retrieved evidence.
Relevance: Answer addresses the query.
Coherence: Output is logically organized.
Toxicity Score: Harmful language measurement.
Latency: Time per request.
Throughput: Requests/tokens per time.
Error Rate: Failed requests percentage.
GPU Utilization: How much GPU is used.

RLHF

RLHF: Aligns LLM using human feedback.
Human Feedback: Human preference labels.
Preference Pair: Two outputs ranked by humans.
Reward Model: Predicts human preference score.
PPO: RL algorithm used in RLHF.
DPO: Directly optimizes preference data.
Reward Hacking: Model exploits reward incorrectly.
Safety Alignment: Makes outputs follow human values.
Constitutional AI: Uses written principles for alignment.
SFT: Supervised fine-tuning on instruction-response pairs.

Software Development

Weight: 24%

Development

Model Deployment: Serve models in production.
API Endpoint: Network interface for model calls.
REST API: HTTP-based service interface.

gRPC: High-performance service protocol.
SDK: Developer tools for using a platform.
Containerization: Package app and dependencies.
Docker: Container runtime.
Kubernetes: Container orchestration system.
Helm Chart: Kubernetes deployment package.
Version Control: Track code/model changes.
Git: Distributed version control tool.
CI/CD: Automated build, test, deployment.
Unit Testing: Tests small code units.
Integration Testing: Tests connected components.
Model Validation Pipeline: Automated model quality checks.
Configuration File: Stores settings/hyperparameters.
YAML/JSON Config: Common configuration formats.
Logging: Records events for debugging/audit.
Asynchronous Programming: Handles many I/O tasks concurrently.
Scalability: Handles growing workload.
Horizontal Scaling: Add more instances.
Vertical Scaling: Add more resources per instance.
Caching: Reuse frequent responses/results.

Frameworks and App Tools

LangChain: Framework for LLM apps.
LangGraph: Orchestrates complex LLM workflows.
ChatNVIDIA: LangChain connector for NVIDIA chat models.
NVIDIAEmbeddings: LangChain connector for NVIDIA embeddings.
Vector Database: Stores/searches embeddings.
FAISS: Vector similarity search library.
Milvus: Vector database.
Pinecone: Managed vector database.
Weaviate: Vector database with semantic search.
ONNX: Open model exchange format.
OpenAI-compatible API: Standard chat/completion API shape.

Monitoring

Performance Monitoring: Tracks latency/throughput/errors.
Model Monitoring: Tracks quality/drift.

Infrastructure Monitoring: Tracks CPU/GPU/memory.

Drift Detection: Finds distribution changes.

Quality Degradation Alert: Warns when model worsens.

Automated Alerts: Real-time issue notifications.

Audit Trail: Record of actions and outputs.

User Feedback Loop: Uses real usage feedback.

Trustworthy AI Principles

Weight: 10%

Ethics

Fairness: Equitable treatment across groups.

Transparency: Decisions are understandable.

Accountability: Clear responsibility for outcomes.

Privacy: Protect sensitive data.

Consent: Respect user data permission.

Beneficence: System should benefit people.

Robustness: Performs reliably under variation.

Explainability: Explains model decisions.

Compliance: Follows laws and regulations.

Model Card: Documents model purpose, limits, risks.

Data Sheet: Documents dataset details.

Bias

AI Bias: Unfair systematic model behavior.

Data Bias: Bias from training data.

Sampling Bias: Unrepresentative data collection.

Label Bias: Biased annotation/ground truth.

Algorithmic Bias: Bias introduced by model/design.

Evaluation Bias: Narrow tests hide failures.

Bias Detection: Measure unfair group differences.

Bias Mitigation: Reduce unfair outcomes.

Fairness Constraints: Enforce fairness during optimization.

Adversarial Debiasing: Train model to reduce bias signals.

Diverse Test Sets: Evaluate across groups/scenarios.

Safety and Guardrails

NeMo Guardrails: Controls safe LLM behavior.

Colang: Language for Guardrails dialogue flows.
Input Rail: Filters user input before LLM.
Output Rail: Filters response after LLM.
Topical Rail: Keeps conversation on allowed topics.
Fact-checking Rail: Verifies claims against sources.
Jailbreak: Prompt that bypasses safety rules.
Prompt Injection: Malicious instruction inside input/context.
Toxicity Detection: Finds harmful language.
PII Protection: Protects personal information.
Differential Privacy: Adds privacy-preserving noise.

Key NVIDIA Technologies

Platforms

NVIDIA NIM: GPU-optimized inference microservice containers.
NVIDIA NeMo: Framework for LLM training/fine-tuning/customization.
TensorRT-LLM: Optimizes LLM inference on NVIDIA GPUs.
Triton Inference Server: Serves models with batching/versioning.
RAPIDS: GPU-accelerated data science libraries.
cuDF: GPU DataFrame library.
cuML: GPU ML algorithms.
cuGraph: GPU graph analytics.
CUDA: NVIDIA GPU programming platform.
cuDNN: GPU primitives for deep neural networks.
NCCL: Multi-GPU communication library.
NGC: NVIDIA catalog for models, containers, SDKs.
BioNeMo: Foundation models for biology/life sciences.
Megatron-LM: Large-scale transformer training framework.
NeMo Curator: Data curation/preprocessing pipeline.
NeMo Guardrails: Safety framework for LLM apps.
NVIDIA AI Endpoints: Hosted NVIDIA model APIs.
DeepStream SDK: Real-time video AI streaming.
Omniverse: 3D simulation/visualization platform.
Omniverse Nucleus: Collaboration/storage for Omniverse assets.

NIM Details

NIM Container: Prebuilt optimized model server.
NGC Distribution: NIM images come from NGC.

OpenAI API Compatibility: Existing clients can switch easily.
REST/gRPC Endpoint: NIM exposes production APIs.
TensorRT-LLM inside NIM: Provides optimized inference.
Triton Backend: Many NIMs use Triton internally.
Hardware Portability: Same container across cloud/on-prem/edge.

NeMo Details

SFT in NeMo: Trains on prompt-response pairs.
RLHF in NeMo: Reward model plus PPO alignment.
LoRA in NeMo: Efficient adapter fine-tuning.
P-Tuning: Learns soft prompt vectors.
Prompt Tuning: Trains prompt embeddings.
Distributed Training: Scales across GPUs/nodes.
Custom Tokenizer: Domain-specific token splitting.
Megatron Integration: Enables tensor/pipeline parallelism.

TensorRT-LLM Details

In-flight Batching: Adds requests into active batches.
Continuous Batching: Keeps GPU busy during serving.
Paged KV Cache: Pages attention cache memory.
Custom CUDA Kernels: Faster low-level GPU operations.
Quantization: Lower precision for speed/memory.
INT8: 8-bit inference precision.
INT4: 4-bit compression precision.
FP8: Low precision floating-point.
AWQ: Activation-aware weight quantization.
GPTQ: Post-training quantization method.
Tensor Parallelism: Splits tensors across GPUs.

Triton Details

Dynamic Batching: Groups requests automatically.
Model Ensemble: Chains models into one pipeline.
Model Versioning: Hosts multiple versions.
Concurrent Execution: Runs multiple models together.
Multi-framework Serving: Supports PyTorch, TF, ONNX, TensorRT.
Model Repository: Directory of deployable models/configs.
A/B Testing: Compare model versions in production.

NGC Details

Model Catalog:	Find pretrained models.
NIM Catalog:	Get production containers.
Helm Charts:	Deploy on Kubernetes.
Version Pinning:	Reproducible container/model versions.
Model Signing:	Verifies container integrity.
Security Scanning:	Checks container vulnerabilities.

Prompt Engineering Mastery

Techniques

Prompt Engineering:	Craft prompts for desired outputs.
Zero-shot Prompting:	No examples provided.
Few-shot Prompting:	Includes examples in prompt.
One-shot Prompting:	Includes one example.
Chain-of-Thought:	Encourages step-by-step reasoning.
Self-consistency:	Sample multiple reason paths; choose consensus.
Tree of Thoughts:	Explore multiple reasoning branches.
Role Prompting:	Assign model a role/persona.
System Prompt:	Sets top-level behavior/instructions.
Delimiters:	Separate instructions from data.
Structured Output Prompting:	Request JSON/tables/schema.
Output Parsing:	Convert model output to usable format.
Prompt Chaining:	Use outputs as later inputs.
Function Calling:	Let LLM call external tools.
Multi-turn Conversation:	Maintain dialogue context.
Domain Adaptation Prompting:	Tailor prompt to domain.
Safety Prompting:	Reduce harmful/off-policy responses.

Generation Controls

Temperature:	Controls randomness.
Low Temperature:	More deterministic output.
High Temperature:	More creative output.
Top-k Sampling:	Sample from k most likely tokens.
Top-p Sampling:	Sample from probability nucleus.
Max Tokens:	Limits output length.
Frequency Penalty:	Discourages repeated tokens.
Repetition Penalty:	Reduces repeated phrases.

- Context Management:** Use context window efficiently.
- Token Efficiency:** Minimize tokens while preserving task.
- Iterative Refinement:** Improve prompt based on output.

Model Evaluation & Testing

Evaluation

- Automated Metrics:** Quantitative model scores.
- Human Evaluation:** Expert/user quality judgment.
- Benchmark Dataset:** Standard comparison dataset.
- GLUE:** NLP understanding benchmark.
- SuperGLUE:** Harder NLP benchmark.
- HELM:** Holistic LLM evaluation benchmark.
- MMLU:** Multi-domain knowledge benchmark.
- Domain-specific Test:** Custom test for target use case.
- Adversarial Test:** Robustness against attacks/edge cases.
- Regression Test:** Ensures behavior does not break.
- Golden Dataset:** Trusted reference test set.

Validation

- Cross-validation:** Reliable performance estimate.
- Hold-out Test:** Final unbiased test set.
- Temporal Validation:** Future data validation.
- Distribution Shift:** Train/test data differ.
- Stress Test:** Performance under heavy load.
- Calibration:** Confidence matches correctness.
- Generalization:** Works on unseen examples.

Monitoring

- Drift Detection:** Input/output distribution changes.
- Performance Tracking:** Continuous metric monitoring.
- Error Analysis:** Investigate failure patterns.
- User Feedback:** Real-world quality signal.
- Automated Alert:** Notifies degradation.
- Production Evaluation:** Measures live system behavior.

Repeated Exam Patterns

Tool-to-Phase Mapping

Data Prep: RAPIDS, cuDF, DALI, NeMo Curator.
Training: NeMo, Megatron-LM, PyTorch, TensorFlow.
Fine-tuning: NeMo, SFT, LoRA, P-Tuning, Prompt Tuning.
Alignment: RLHF, reward model, PPO, DPO.
Optimization: TensorRT-LLM, quantization, batching, KV cache.
Serving: NIM, Triton, REST/gRPC.
App Development: LangChain, LangGraph, NVIDIA AI Endpoints.
Retrieval: Embeddings, vector DB, FAISS, Milvus.
Safety: NeMo Guardrails, Colang, rails.
Catalog: NGC.
Monitoring: Logs, drift, latency, throughput, alerts.

RAG Pipeline

User Query: User asks a question.
Query Embedding: Convert query to vector.
Vector Database: Stores searchable document vectors.
ANN Search: Fast approximate nearest neighbor retrieval.
Retriever: Finds relevant chunks.
Re-ranker: Reorders retrieved chunks by relevance.
Context Injection: Add chunks into prompt.
LLM Generation: Produce grounded answer.
Citation/Source Check: Verify answer support.
RAG Benefit: Reduces hallucination and adds fresh knowledge.
RAG Risk: Bad retrieval causes bad answers.

Fine-tuning Methods

Full Fine-tuning: Updates all model weights.
PEFT: Updates small trainable components.
LoRA: Low-rank trainable adapters.
QLoRA: Quantized base model plus LoRA.
Adapter Tuning: Small modules inserted in layers.
Prompt Tuning: Trains soft prompt tokens.
P-Tuning: Learns continuous prompt vectors.
Instruction Tuning: Teaches instruction following.
Domain Adaptation: Adapts model to specialized domain.
Catastrophic Forgetting: New training erases old skills.
Gradient Checkpointing: Saves memory by recomputing activations.

Mixed Precision Training: Uses FP16/BF16 for speed/memory.

Optimization and Deployment

Quantization: Reduces precision for efficient inference.

Pruning: Removes less important weights.

Distillation: Smaller model learns from larger model.

Batching: Process requests together.

Dynamic Batching: Triton groups concurrent requests.

Continuous/In-flight Batching: Adds requests mid-generation.

Speculative Decoding: Smaller draft model speeds generation.

KV Cache: Speeds autoregressive decoding.

Caching: Stores common responses/results.

Autoscaling: Adjusts replicas to demand.

Cold Start: Delay when starting service/model.

Edge Deployment: Run model near device/user.

Model Registry: Tracks model versions/artifacts.

Canary Release: Small rollout before full release.

Container Orchestration: Manage containers at scale.

Trust and Risk

Hallucination: Confident but false output.

Grounding: Connect output to evidence.

Jailbreak: Bypass safety guardrails.

Prompt Injection: Malicious prompt hidden in input.

Data Leakage: Sensitive data exposed.

Copyright Risk: Training data may violate rights.

Privacy Risk: User data mishandled.

Bias Risk: Unequal performance across groups.

Toxic Output: Harmful or offensive response.

Robustness Risk: Fails under unusual inputs.

Accountability Gap: Unclear responsibility for harm.

Data Types and Preprocessing

Image Data: Resize, normalize, augment.

Text Data: Tokenize, pad, embed.

Audio Data: Spectrograms, normalization, augmentation.

Tabular Data: Scale numbers, encode categories.

Time-series Data: Normalize, window, forecast.

Image Rotation: Augments image angle.

Image Flipping: Mirrors images.

Random Crop: Varies image view.

Color Jitter: Changes brightness/contrast/color.

Noise Injection: Adds random noise for robustness.

Synonym Replacement: Text augmentation by synonyms.

Back Translation: Translate out and back to paraphrase.

Random Insertion/Deletion: Adds/removes words.

Time Shift: Audio time augmentation.

Pitch Scaling: Audio pitch augmentation.

Common Confusions

NIM vs Triton: NIM is packaged model microservice; Triton is serving server.

NIM vs NGC: NIM runs models; NGC stores/distributes containers/models.

NeMo vs NIM: NeMo trains/customizes; NIM deploys/serves.

TensorRT-LLM vs Triton: TensorRT-LLM optimizes LLMs; Triton serves models.

Guardrails vs RLHF: Guardrails control runtime behavior; RLHF changes model alignment.

RAPIDS vs Pandas: RAPIDS/cuDF runs DataFrames on GPU; pandas runs on CPU.

RAG vs Fine-tuning: RAG retrieves knowledge; fine-tuning changes weights.

LoRA vs Full Fine-tuning: LoRA updates adapters; full updates all weights.

BLEU vs ROUGE: BLEU often translation; ROUGE often summarization.

Precision vs Recall: Precision avoids false positives; recall avoids false negatives.

Latency vs Throughput: Latency is delay; throughput is volume.

Encoder vs Decoder: Encoder understands; decoder generates.

Self-attention vs Cross-attention: Same sequence vs encoder-decoder attention.

AI Subject Terms Glossary

Core AI Terms

Agent: System that acts toward a goal.

AI Pipeline: Steps from data to deployed model.

Algorithm: Procedure for solving a task.

Baseline: Simple reference model/performance.

Checkpoint: Saved model state.

Class Imbalance: Unequal class distribution.

Data Leakage: Test information enters training.

Decision Boundary: Separates predicted classes.

Embedding Space: Vector space of learned meanings.

Feature: Input variable used by model.
Generalization: Performs well on unseen data.
Ground Truth: Correct reference answer/label.
Heuristic: Rule-of-thumb method.
Label: Correct output for supervised training.
Latent Space: Compressed hidden representation.
Model Architecture: Structure of model layers/components.
Objective Function: Function optimized during training.
Pretraining: Training on broad large data first.
Prompt: Input instructions sent to LLM.
Representation: Learned encoding of data.
Sampling: Choosing output tokens probabilistically.
Training Corpus: Dataset used for training.
Validation Set: Data used for tuning decisions.

Neural Network Terms

Activation Function: Adds nonlinearity to networks.
ReLU: Outputs $\max(0, x)$.
Sigmoid: Maps values to 0-1.
Tanh: Maps values to -1 to 1.
GELU: Smooth activation common in transformers.
Dead ReLU: ReLU stuck outputting zero.
Gradient: Direction/size of parameter update.
Vanishing Gradient: Gradients become too small.
Exploding Gradient: Gradients become too large.
Optimizer: Updates model weights.
AdamW: Adaptive optimizer with weight decay.
SGD: Basic gradient descent optimizer.
Weight Decay: Regularization penalizing large weights.
Batch Normalization: Normalizes layer inputs.
Layer Normalization: Normalizes features per sample.
Residual Block: Layer with skip connection.
Parameter Count: Number of learned weights.
Compute: Processing needed for training/inference.
FLOPs: Floating-point operations.
VRAM: GPU memory.

NLP and LLM Terms

NLP: AI for human language.
Token: Basic text unit processed by model.
Tokenizer: Converts text into tokens.
Vocabulary: Set of tokenizer tokens.
Subword Tokenization: Splits words into pieces.
BPE: Byte Pair Encoding tokenizer method.
WordPiece: Subword tokenizer used by BERT.
SentencePiece: Language-independent tokenizer.
Padding Token: Filler token for equal lengths.
Attention Mask: Marks real tokens vs padding.
Autoregressive Model: Predicts next token from previous tokens.
Masked Language Modeling: Predicts hidden tokens.
Next Sentence Prediction: Predicts sentence continuity.
Causal Language Modeling: Next-token prediction.
Bidirectional Attention: Looks left and right.
Unidirectional Attention: Looks only previous tokens.
Instruction Following: Model obeys user task instructions.
In-context Learning: Learns from prompt examples.
Context Length: Maximum tokens accepted.
Stop Sequence: Text that ends generation.
Logit: Raw score before probability.
Softmax: Converts scores to probabilities.
Decoding: Converts model probabilities to text.
Greedy Decoding: Always picks most likely token.
Beam Search: Keeps multiple candidate sequences.
Nucleus Sampling: Top-p sampling.

Generative AI Terms

Generative Model: Creates new data.
Discriminative Model: Predicts labels/classes.
Generator: GAN network that creates samples.
Discriminator: GAN network that judges samples.
Forward Diffusion: Adds noise to data.
Reverse Diffusion: Removes noise to generate.
Denoising: Removing noise from data.
Latent Diffusion: Diffusion in compressed space.
Multimodal Model: Handles multiple input types.

Image Captioning: Generates text description of image.
Text-to-Image: Generates image from prompt.
Image-to-Text: Generates text from image.
Speech-to-Text: Converts audio to text.
Text-to-Speech: Converts text to audio.
Synthetic Data: Artificially generated training data.
Hallucination: Plausible but incorrect generation.
Grounded Generation: Generation supported by evidence.

Evaluation Terms

Benchmark: Standard test for comparison.
Metric: Numeric performance measure.
Loss: Training error signal.
Train Loss: Error on training data.
Validation Loss: Error on validation data.
Test Accuracy: Final held-out accuracy.
False Positive: Incorrect positive prediction.
False Negative: Incorrect negative prediction.
True Positive: Correct positive prediction.
True Negative: Correct negative prediction.
ROC Curve: TPR/FPR curve across thresholds.
Confidence Score: Model certainty estimate.
Calibration Error: Confidence mismatch.
Inter-rater Agreement: Human judge consistency.
Regression Metric: Score for numeric prediction.
Classification Metric: Score for class prediction.
Retrieval Metric: Score for search quality.
Recall@k: Relevant item appears in top k.
MRR: Mean reciprocal rank.
NDCG: Ranking quality metric.

Data Science Terms

Grid Search: Tests all parameter combinations.
Randomized Search: Samples parameter combinations.
Hyperparameter Tuning: Finds best config.
Cross-validation Fold: One split in k-fold testing.
ARIMA: Time-series forecasting model.

Autocorrelation:	Correlation with past values.
Partial Autocorrelation:	Direct lag correlation.
Stationarity:	Stable mean/variance over time.
Seasonality:	Repeating time pattern.
Trend:	Long-term direction.
Intercept:	Baseline regression value.
Coefficient:	Feature effect in regression.
Residual:	Prediction error.
Outlier:	Unusual data point.
Missing Value:	Empty/unknown feature value.
Imputation:	Fill missing values.
One-hot Encoding:	Converts categories to binary columns.

NVIDIA and Production Terms

MIG:	Splits one GPU into isolated instances.
Time-sharing:	Multiple workloads share GPU time.
Model Signing:	Cryptographic verification of model/container.
Vulnerability Scanning:	Checks security weaknesses.
Version Pinning:	Locks exact model/container version.
Helm:	Kubernetes package manager.
Load Balancing:	Distributes traffic across servers.
SLA:	Required service reliability/performance.
Observability:	Logs, metrics, traces for systems.
Telemetry:	Collected runtime measurements.
Throughput Bottleneck:	Component limiting system speed.
CPU-GPU Transfer:	Data movement between CPU and GPU.
Zero-copy Interop:	Avoids unnecessary data copying.

Python Tools and Libraries

Core Python ML Stack

Python:	Main language for AI/ML workflows.
NumPy:	Fast numerical arrays and math.
Pandas:	CPU DataFrame data analysis.
Polars:	Fast DataFrame processing library.
SciPy:	Scientific computing and optimization.
scikit-learn:	Classical ML and metrics toolkit.
Jupyter Notebook:	Interactive coding and exploration.

Matplotlib: Basic Python plotting.

Seaborn: Statistical visualization on Matplotlib.

Plotly: Interactive charts and dashboards.

Statsmodels: Statistical models and time-series tools.

NLTK: Traditional NLP toolkit.

spaCy: Production NLP pipeline library.

Deep Learning Libraries

PyTorch: Flexible deep learning framework.

torch: PyTorch tensor and neural network package.

torch.nn: PyTorch neural network layers.

torch.optim: PyTorch optimizers.

torchvision: PyTorch image models/transforms.

torchaudio: PyTorch audio processing.

TensorFlow: Production deep learning framework.

Keras: High-level TensorFlow model API.

TensorBoard: Training visualization dashboard.

JAX: High-performance array/autodiff framework.

Autograd: Automatic gradient computation.

CUDA Tensor: Tensor stored on NVIDIA GPU.

Hugging Face Ecosystem

Transformers: Pretrained transformer models library.

AutoModel: Loads model architecture automatically.

AutoTokenizer: Loads matching tokenizer automatically.

Datasets: Dataset loading and preprocessing library.

Tokenizers: Fast tokenizer implementations.

Accelerate: Simplifies distributed/mixed-precision training.

PEFT: Parameter-efficient fine-tuning library.

TRL: RLHF/DPO training utilities.

Safetensors: Safe fast model weight format.

Model Hub: Repository for pretrained models/datasets.

Pipeline API: Quick task-based inference wrapper.

LLM App Development

LangChain: Builds LLM apps and chains.

LangGraph: Stateful graph workflows for agents.

LlamaIndex: Data connectors and RAG framework.

OpenAI Python SDK: Calls OpenAI-compatible APIs.
NVIDIA AI Endpoints: Hosted NVIDIA model APIs.
ChatNVIDIA: LangChain chat connector for NVIDIA.
NVIDIAEmbeddings: LangChain embedding connector.
FAISS Python: Local vector similarity search.
Milvus SDK: Connects Python to Milvus vector DB.
Pinecone SDK: Connects Python to Pinecone.
Weaviate Client: Connects Python to Weaviate.
Chroma: Lightweight local vector database.
SentenceTransformers: Embedding models for semantic search.
tiktoken: Token counting/tokenization utility.

NVIDIA Python and GPU Tools

RAPIDS: GPU data science Python ecosystem.
cuDF: GPU DataFrame library.
cudf.pandas: GPU acceleration for pandas code.
cuML: GPU machine learning algorithms.
cuGraph: GPU graph analytics.
Dask-cuDF: Distributed GPU DataFrames.
CuPy: NumPy-like arrays on GPU.
Numba: JIT compiler for Python/GPU kernels.
NVIDIA DALI: GPU data loading and preprocessing.
NVTabular: GPU tabular feature engineering.
NeMo Toolkit: Python framework for LLM customization.
NeMo Curator: Dataset curation and filtering.
NeMo Guardrails: Python safety/rail framework.
TensorRT Python API: Build optimized inference engines.
TensorRT-LLM: Optimized LLM inference library.
Triton Client: Python client for Triton inference.
NVIDIA NIM Client Pattern: Call NIM via REST/gRPC.
NCCL Bindings: Multi-GPU communication support.

Data Processing Libraries

PyArrow: Columnar data and Parquet support.
Parquet: Efficient columnar storage format.
JSONL: One JSON object per line.
CSV: Simple tabular text format.

BeautifulSoup: HTML parsing.

Requests: HTTP API calls.

aiohttp: Async HTTP client/server.

Regex: Pattern matching for text cleaning.

Unidecode: Text normalization utility.

pydantic: Data validation and schemas.

Serving and API Libraries

FastAPI: Modern Python REST API framework.

Flask: Lightweight Python web framework.

Uvicorn: ASGI server for FastAPI.

Gunicorn: Production Python web server.

gRPC Python: High-performance RPC services.

Streamlit: Quick Python data apps.

Gradio: Quick ML demos/interfaces.

ONNX Runtime: Runs ONNX models.

Triton Inference Server: Production model serving.

Docker SDK: Control Docker from Python.

Kubernetes Client: Manage Kubernetes from Python.

Testing and Quality Tools

pytest: Python testing framework.

unittest: Built-in Python test framework.

mypy: Static type checking.

ruff: Fast Python linter/formatter.

black: Opinionated Python formatter.

pre-commit: Runs checks before commits.

tox: Tests across environments.

coverage.py: Measures test coverage.

Great Expectations: Data quality validation.

Evidently: ML drift and monitoring reports.

Experiment Tracking and MLOps

MLflow: Tracks experiments and models.

Weights & Biases: Experiment tracking dashboard.

DVC: Data/model version control.

Hydra: Configuration management.

OmegaConf: Hierarchical config library.

Airflow: Workflow orchestration.

Prefect: Python workflow orchestration.

Ray: Distributed Python computing.

Ray Serve: Scalable model serving.

BentoML: Package and serve ML models.

Python Concepts for Certification

Virtual Environment: Isolates project dependencies.

pip: Python package installer.

conda: Environment and package manager.

requirements.txt: Lists pip dependencies.

pyproject.toml: Modern Python project config.

Asyncio: Python async concurrency.

Coroutine: Async function execution unit.

Type Hint: Optional Python type annotation.

Exception Handling: Handles runtime errors.

Serialization: Save/load data structures.

Pickle: Python object serialization.

Environment Variable: Runtime configuration value.

Logging Module: Built-in structured logging.

Exam Decision Tree

NVIDIA Tool Choice

Need deploy model fast: Choose NVIDIA NIM.

Need OpenAI-compatible endpoint: Choose NVIDIA NIM.

Need private/on-prem inference: Choose NVIDIA NIM.

Need optimize LLM inference: Choose TensorRT-LLM.

Need lower latency: Choose TensorRT-LLM or batching.

Need higher throughput: Choose TensorRT-LLM, Triton, batching.

Need serve many frameworks: Choose Triton.

Need model versioning/A-B tests: Choose Triton.

Need chain models server-side: Choose Triton ensemble.

Need train or fine-tune LLM: Choose NeMo.

Need RLHF/SFT/LoRA: Choose NeMo.

Need curate training data: Choose NeMo Curator.

Need GPU DataFrame processing: Choose RAPIDS/cuDF.

Need GPU ML algorithms: Choose cuML.

Need GPU graph analytics: Choose cuGraph.

Need safety/policy controls: Choose NeMo Guardrails.

Need block jailbreaks/injections: Choose Guardrails rails.

Need find NVIDIA containers/models: Choose NGC.

Need biology foundation models: Choose BioNeMo.

Need video AI streaming: Choose DeepStream.

Need GPU preprocessing: Choose DALI.

Model Method Choice

Need external/current knowledge: Use RAG.

Need domain behavior/style: Use fine-tuning.

Need cheap fine-tuning: Use PEFT/LoRA.

Need align to preferences: Use RLHF/DPO.

Need reduce model size: Use distillation/pruning/quantization.

Need faster decoding: Use KV cache/speculative decoding.

Need reduce memory: Use quantization/LoRA/checkpointing.

Need compare live versions: Use A/B testing.

Need unbiased final score: Use held-out test set.

Need robust estimate: Use cross-validation.

Common Exam Traps

RAG vs Fine-tuning: RAG adds knowledge; fine-tuning changes behavior/weights.

NIM vs NGC: NIM runs models; NGC stores models/containers.

NIM vs Triton: NIM is packaged service; Triton is serving infrastructure.

NeMo vs NIM: NeMo trains/customizes; NIM deploys.

TensorRT-LLM vs Triton: TensorRT optimizes; Triton serves.

Guardrails vs RLHF: Guardrails runtime controls; RLHF training alignment.

cuDF vs Pandas: cuDF GPU; pandas CPU.

cuML vs scikit-learn: cuML GPU ML; scikit-learn CPU ML.

Latency vs Throughput: Latency delay; throughput volume.

Precision vs Recall: Precision reduces false positives; recall reduces false negatives.

F1 vs Accuracy: F1 better for imbalance; accuracy can mislead.

BLEU vs ROUGE: BLEU translation; ROUGE summarization.

Perplexity vs Accuracy: Perplexity language uncertainty; accuracy exact correctness.

Encoder-only vs Decoder-only: Encoder understands; decoder generates.

BERT vs GPT: BERT bidirectional encoder; GPT autoregressive decoder.

Self-attention vs Cross-attention: Same sequence vs encoder-decoder link.

Top-k vs Top-p: Fixed count vs probability mass.

Temperature 0 vs High: Deterministic vs creative.

Quantization vs Pruning: Lower precision vs remove weights.

Data Parallelism vs Tensor Parallelism: Split data vs split model tensors.

Pipeline Parallelism vs Tensor Parallelism: Split layers vs split tensor operations.

Validation Set vs Test Set: Tune on validation; final score on test.

Data Bias vs Algorithmic Bias: Bad data vs biased model/design.

Prompt Injection vs Jailbreak: Hidden malicious instruction vs bypass attempt.

Groundedness vs Relevance: Supported by sources vs answers question.

Formula and Metric Mini Sheet

Classification

Accuracy: Correct predictions / all predictions.

Precision: $TP / (TP + FP)$.

Recall: $TP / (TP + FN)$.

F1: $2 * \text{precision} * \text{recall} / (\text{precision} + \text{recall})$.

False Positive: Predicted positive, actually negative.

False Negative: Predicted negative, actually positive.

Confusion Matrix: Table of TP, FP, FN, TN.

AUC-ROC: Ranking quality across thresholds.

Regression

MAE: Average absolute error.

MSE: Average squared error.

RMSE: Square root of MSE.

R-squared: Proportion of variance explained.

NLP and LLM

Perplexity: Lower means better next-token prediction.

Cross-entropy: Penalizes wrong probability predictions.

BLEU: N-gram precision for translation.

ROUGE: N-gram recall for summarization.

METEOR: Translation metric with synonym/stem matching.

Groundedness: Claims supported by context.

Faithfulness: Output does not contradict source.

Toxicity: Harmful/offensive content score.

Retrieval and RAG

Cosine Similarity: Measures vector angle similarity.
Recall@k: Relevant item appears in top k.
Precision@k: Relevant fraction in top k.
MRR: Average reciprocal rank of first relevant result.
NDCG: Ranking quality with position weighting.

Production

Latency: Time for one request.
Throughput: Requests/tokens per second.
Error Rate: Failed requests / total requests.
GPU Utilization: Percent GPU busy.
Memory Footprint: RAM/VRAM required.
Tokens/sec: Generation speed.

Scenario Keywords

NVIDIA Tools

"OpenAI-compatible API": NIM.
"Pre-packaged optimized container": NIM.
"Run model locally/private cloud": NIM.
"Pull from nvcr.io": NGC/NIM.
"Catalog of models/containers": NGC.
"Helm chart": NGC/Kubernetes deployment.
"Fine-tune LLM": NeMo.
"SFT/RLHF/LoRA/P-Tuning": NeMo.
"Reward model/PPO": RLHF with NeMo.
"Data curation/filtering": NeMo Curator.
"Block unsafe topics": NeMo Guardrails.
"Colang": NeMo Guardrails.
"Input/output rails": NeMo Guardrails.
"Optimize LLM latency": TensorRT-LLM.
"Paged KV cache": TensorRT-LLM.
"INT4/INT8/FP8": TensorRT-LLM quantization.
"Dynamic batching": Triton.
"Model ensemble": Triton.
"Model versioning": Triton.

"Serve PyTorch/TensorFlow/ONNX together": Triton.

"GPU pandas": cuDF.

"GPU data science": RAPIDS.

"GPU clustering/PCA/nearest neighbors": cuML.

"GPU graph analytics": cuGraph.

Model and Data Methods

"Needs latest/private documents": RAG.

"Reduce hallucinations with sources": RAG.

"Vector database": Embeddings/RAG.

"Semantic search": Embeddings/vector DB.

"Small GPU memory for tuning": LoRA/QLoRA.

"Train only few parameters": PEFT.

"Human preferences": RLHF/DPO.

"Pairwise comparisons": Reward model/DPO.

"Reduce model size": Distillation/pruning/quantization.

"Faster inference": Quantization, TensorRT-LLM, batching.

"Multiple GPUs too large model": Tensor parallelism.

"Split layers across GPUs": Pipeline parallelism.

"Same model copy on each GPU": Data parallelism.

Evaluation and Safety

"Imbalanced classes": F1/precision/recall.

"False positives costly": Precision.

"False negatives costly": Recall.

"Translation quality": BLEU/METEOR.

"Summarization quality": ROUGE/human eval.

"Language model fluency": Perplexity.

"Open-ended answers": Human eval/LLM-as-judge.

"Production comparison": A/B testing.

"Final unbiased score": Hold-out test set.

"Distribution changes": Drift detection.

"Harmful prompts": Red teaming/adversarial testing.

"Sensitive user data": Privacy/PII controls.

"Demographic fairness": Bias/fairness evaluation.